

Win 1–2 labs deep. Ship samples that earn pilots. Staff the bottlenecks.

Outsized growth concentrates in a few foundation labs. Enterprise is a controlled second lane built with GLG’s core commercial relationships.

Patrick Severs · General Manager, AI / Human Data · Final-round case · August 2026

Thesis Professional-domain human data beyond trust and safety.	Motion Relationships → rooms → samples → pilots → expand by team.	Build Mobilize internal bench. Add Solutions, invite operator, and data science. Protect Delivery.
--	---	--

If it does not create a sample, a paid pilot path, or remove a hard bottleneck in 90 days, it waits.

Become a core lab training partner by productizing professional judgment.

MSA paper and trust-and-safety work are the start. The gap is proactive GTM, sample-led solutioning, and a quality system labs expand.

What I believe

- Labs pay for accepted professional signal, not expert access alone
- Finance, legal, healthcare, and business ops are GLG's natural wedges
- \$100M-scale growth concentrates in 1-2 labs that explode by team and wedge
- Enterprise is real but second; open it with core commercial, don't wait

What I will do

- Open Wave-1 rooms where I already have access
- Ship account-specific samples buyers can evaluate
- Convert accepted samples into paid pilots and POs
- Invest only where a measured bottleneck blocks that loop

North star: human data operating layer for professional expertise. Quality, urgency, agility.

Sell the work labs already buy in professional domains.

Public pattern: experienced practitioners author real work products, other practitioners judge them. That is the GLG wedge. Seen across GDPval-style and post-training evaluation work.

Finance + risk Labs want: credit memos, investment memos, risk writeups, financial reasoning rankings, grounded research answers. GLG builds: CFA/credit/IB-style cohorts grading model drafts and writing gold answers.	Legal + compliance Labs want: brief quality, clause risk, policy exceptions, jurisdiction-aware critiques, escalation calls. GLG builds: practice-area panels for preferred drafts, failure labels, and release gates.	Healthcare + LS Labs want: care-plan quality, clinical reasoning, safety omissions, protocol adherence. GLG builds: clinician/regulatory panels for gold outputs and critical-error flags.	Business ops Labs want: realistic enterprise workflows, decision usefulness, tool-use judgment where humans score outcomes. GLG builds: operator panels for tasks, edge cases, and trace review without owning the agent stack.
--	--	--	---

① Tasks Solutions builds: prompt, source pack, output format, edge cases. Must look like real work.	→ ② Gold answers Experts + Delivery: decision-grade reference work and rationale a senior practitioner accepts.	→ ③ Judgment Delivery runs: blind rank/critique/correct, agreement check, adjudication on hard cases.	→ ④ Holdout Solutions + Delivery: unseen set for model comparison or release. Never mixed into training.
--	---	---	--

Buyers ask for sample data. That sample decides whether you get a pilot.

They are evaluating three things: do you understand the domain, can you deliver accurate accepted work, and can you turn it around without drama.

TASK_ID	DOMAIN	PROMPT / ASK	GOLD_ANSWER (EXCERPT)	MODEL_RANK	FAILURE_LABEL	RATIONALE	SPLIT
FIN-014	Credit	Draft downside risks for a mid-market borrower using the attached 10-K excerpts.	Names leverage spike, customer concentration, covenant headroom; flags missing liquidity stress.	B > A > C	material_omission	A ignores supplier concentration; C invents a facility not in sources.	train
FIN-015	Credit	Same setup, unseen borrower pack.	Held-out gold memo + critical-error rules.	—	—	Used only to measure lift after pilot data lands.	holdout
LEG-006	Compliance	Review model summary of a supplier contract for missing escalation clauses.	Marks indemnity gap + audit-rights weakness; prefers revised clause set.	A > C > B	unsupported_claim	B invents governing law; C underweights termination risk.	train

What the sample proves

- We can match the right practitioners
- We can write tasks and rubrics labs recognize
- We can deliver structured, rights-aware data objects

How it earns the pilot

- Named acceptance owner reviews the batch
- High accepted-first-time / low rework = trust
- Pilotful signal → paid pilot SOW with expansion decision

~80% labs. Wave 1 gets the samples. Named people open the rooms.

All current AI-training revenue sits under MSAs. Paper is access. Samples and POs are the business.

Where time goes

Wave 1: Amazon AGI · Microsoft AI · OpenAI · Mistral

First samples land here. Expand by team and wedge inside the winner.

Wave 2: DeepMind · xAI · Cohere

Path-building. Sample only after a live need is confirmed.

Enterprise ~20%: one design partner from existing GLG corporate relationships.

Priority unlock paths

- **Annika Huston** · Amazon AGI Post-Training Data & Evals
- **Nguyen Bach** · Microsoft AI / data acquisition
- **Wyatt Thompson** · OpenAI Human Data
- **Anuj Nair** · OpenAI technical route
- **Darren Cosker** · Microsoft technical validator
- **Prateek / Howard → Eujeong / Kahina** · Mistral
- **Georgi Stoyanov** · DeepMind data ops
- **Phil Holland** · xAI post-training

Full top-10 stakeholder table with first asks is in the appendix.

Win with relationships and small rooms. Not an SDR machine.

No giant booths. No carpet-bombing the conference floor. Build environments people want to join, then convert.

1 · PARTNER ROOMS

Open the right doors

AI Circle. Influential ambassadors. Highly connected individuals. Selective partnerships with human-data companies that cover different domains than GLG.

2 · SHOW UP WHERE THEY ARE

ICML · NeurIPS · COLM

Be present for credibility and listening. Do not carpet-bomb the venue. Set up the adjacent room people actually want to enter.

3 · INTIMATE GTM

Dinners · working sessions · individual attention

SF / NYC / London. ~5–12 people. Me plus a Solutions or domain partner. Create the gathering. Build the relationship.

4 · CONVERT

Brief → sample → pilot

Every strong room ends in a problem brief or a clear no. Accepted sample earns the paid pilot. Expand by team, then feed the next room.

↻ expand the account · open the next room · repeat

Transparent invite. Low ask. Clear path to a sample.

These are relationship openers, not sequences from an SDR playbook.

Email 1

To: Annika Huston · Amazon AGI

Subject: Small dinner on what good human data looks like

Annika,

I'm building GLG's AI training and evals business. I don't want to start with a polished pitch.

I'd like to bring a small group who run post-training, evals, and human data around one table in San Francisco, paired with two or three GLG experts who understand real business workflows.

Question: where does quality break today, and where could experienced professional judgment produce a better signal?

No panel. No vendor deck. If useful, we follow with a small sample.

Open to joining and helping shape the room?

Patrick

Email 2

To: Nguyen Bach · Microsoft AI

Subject: Working session on professional-domain human data

Nguyen,

I'm leading GLG's AI training build. GLG already supports high-stakes specialist work with labs and has deep professional networks in finance, legal, healthcare, and operations.

I'd value a short working session on where professional-domain data and evals are most constrained for your teams outside coding.

If useful, I can bring one Solutions partner and a domain expert, then follow with a tightly scoped sample against a live capability gap.

Worth 30 minutes this month?

Patrick

These are the conditions that keep the business from breaking.

The operating system we must put in place so samples convert and quality holds.

Must be in place

- **Named technical buyer + acceptance owner** before custom sample work
- **Paper-to-people map** so MSA is never mistaken for funded demand
- **Calibrated cohorts before volume** with backup experts ready
- **First-pass acceptance gate** before any scale-up
- **Separate Solutions and Delivery ownership** so one person is not doing both
- **Security / tooling path running in parallel** with the sample
- **Working-capital plan** before concurrent pilots pay experts ahead of collection


Most important to get done

- 1 Open Wave-1 rooms and leave with sample briefs
- 2 Ship accepted samples that earn pilot decisions
- 3 Seat Solutions and protect Delivery quality
- 4 Prove one lab expansion path by team and wedge
- 5 Stand up one enterprise design partner with core commercial

If these five move, the rest of the org can follow the work.

Small core. Clear owners. Mobilize delivery. Staff the rooms.

Mobilize the internal bench first. Augment with Solutions and data science. Add an early GTM operator so invite logistics do not bury the GM.



Partner with core commercial. Find one design partner in the install base.

Enterprise human data is early and has not been a top priority for teams that own those relationships.
So we build the motion together.

Offer

Gold Set + Challenge Sprint

One live regulated workflow. Scenarios, gold answers, rubrics, corrections, held-out challenge. Client keeps model and stack.

How we open

- Sit with core GLG commercial / account leaders
- Map existing corporate clients with live AI workflows
- Pick one design partner, not a logo hunt

90-day bar

- Shared shortlist with commercial
- A few qualified discovery conversations
- One design-partner brief or paid signal phase

Same quality system as labs. Different packaging. Do not steal sample capacity from a live lab pilot.

Concrete outputs. Open the rooms. Ship the samples.

Timing on specific people depends on calendars. The motion does not.

DAYS 0-30

Map + open

- MSA → SOW/PO → buyer map
- 3 Wave-1 account plans
- Name mobilizable internal bench
- Start outreach / dinners / working sessions with Wave-1 paths
- Working session with core commercial on enterprise install base
- Open SA + data science search

DAYS 31-60

Build proof

- Ship account-specific samples where briefs are ready
- Calibrate 2 expert cohorts
- Run security/tooling path with those samples
- Advance enterprise shortlist with commercial
- SA / Delivery finalists

DAYS 61-90

Convert

- Drive ≥1 paid pilot path
- Expansion / PO conversation on the winner
- One enterprise design-partner brief or paid phase if ready
- SA seated or contracted; Delivery owner named
- Reforecast from observed conversion

Days 91-180: double down on the winning lab. Add people against signed load. Keep enterprise to one design partner until the factory is stable.

Go big with 1–2 labs. Do not build a balanced logo fantasy.

Concentration is the honest bet. Underwrite it and manage the risk in the open.

How we get there

- 1–2 anchor labs become the majority of revenue
- Expand by team / wedge / program inside those accounts
- A few secondary lab programs for diversification
- Enterprise as a smaller recurring lane after the factory works

What has to be true

- At least one lab has expansion surface beyond T&S
- Samples convert to pilots and quarterly POs
- Accepted yield holds as volume rises
- Working capital and concentration limits are explicit

Public ceiling signal: one Innodata account reached roughly \$110.5M annualized in a year, with heavy concentration. Possible. Not automatic. Source: SEC filings.

No fake ramp. Day 90 evidence decides whether to double down on the concentration bet.

Clear ownership. Fast samples. No shiny objects.

I set strategy, open rooms, and bring executors. I am not the technical operator.

I own

Priority, rooms, scarce sample capacity, commercial next step.

Solutions owns

Need → tasks, rubrics, sample design.

Delivery owns

Acceptance, calibration, stop-the-line.

Hard no's

- No commodity annotation / no platform before wedge proof
- No multi-logo enterprise hunt
- No untargeted event spend without a sample path

If this lands

- Meet core commercial / product on enterprise design-partner path
- Start paper-to-people audit day one
- Open Wave-1 relationship motion immediately

If it does not create a sample, a paid pilot path, or remove a hard bottleneck in 90 days, it waits.